

project2

April 8, 2025

1 Machine Learning in Python - Project 2

Due Friday, Apr 11th by 4 pm.

Include contributors names in notebook metadata or here

1.1 Setup

Install any packages here, define any functions if needed, and load data

```
[1]: # Add any additional libraries or submodules below

# Data libraries
import pandas as pd
import numpy as np

# Plotting libraries
import matplotlib.pyplot as plt
import seaborn as sns

# Plotting defaults
plt.rcParams['figure.figsize'] = (8,5)
plt.rcParams['figure.dpi'] = 80

# sklearn modules
import sklearn
from sklearn.model_selection import train_test_split
from sklearn.impute import SimpleImputer
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.metrics import (
    accuracy_score, precision_score, recall_score, f1_score,
    ↪precision_recall_curve,
    roc_auc_score, confusion_matrix, ConfusionMatrixDisplay, RocCurveDisplay,
    ↪classification_report
)
```

```

from imblearn.pipeline import Pipeline as imPipeline
from imblearn.over_sampling import SMOTE
import xgboost as xgb
from xgboost import XGBClassifier
from sklearn.model_selection import RandomizedSearchCV, StratifiedKFold
import shap
from scipy.stats import skew

```

```

[2]: # Load data in easyshare.csv
d = pd.read_csv("freddiemac.csv", dtype={
    'cd_msa': str,
    'zipcode': str,
    'id_loan_rr': str,
    'rr_ind': str})
d.head()

```

```

[2]:    fico  dt_first_pi  flag_fthb  dt_matr  cd_msa  mi_pct  cnt_units  occpy_sts  \
0    809         201705         N   204704     NaN        0           1          P
1    702         201703         N   203202     NaN        0           1          P
2    792         201703         N   204702     NaN        0           1          S
3    776         201703         N   204702     NaN        0           1          S
4    790         201703         N   204702  41620        0           1          I

```

```

    cltv  dti  ...  seller_name  servicer_name  flag_sc  \
0    75   38  ...  Other sellers  SPECIALIZED LOAN SERVICING LLC    NaN
1    80   36  ...  Other sellers          Other servicers    NaN
2    60   36  ...  Other sellers          Other servicers    NaN
3    80   18  ...  Other sellers          Other servicers    NaN
4    75   42  ...  Other sellers          PNC BANK, NA      NaN

```

```

    id_loan_rr  program_ind  rr_ind  property_val  io_ind  mi_cancel_ind  loan_status
0         NaN           9     NaN           2     N           7      prepaid
1         NaN           9     NaN           2     N           7      active
2         NaN           9     NaN           2     N           7      prepaid
3         NaN           9     NaN           2     N           7      prepaid
4         NaN           9     NaN           2     N           7      active

```

[5 rows x 33 columns]

2 Introduction

This section should include a brief introduction to the task and the data (assume this is a report you are delivering to a professional body (e.g. Freddie Mac).

Briefly outline the approaches being used and the conclusions that you are able to draw.

3 Exploratory Data Analysis and Feature Engineering

Include a detailed discussion of the data with a particular emphasis on the features of the data that are relevant for the subsequent modeling. Including visualizations of the data is strongly encouraged - all code and plots must also be described in the write up. Think carefully about whether each plot needs to be included in your final draft and the appropriate type of plot and summary for each variable type - your report should include figures but they should be as focused and impactful as possible.

You should also split your data into training and testing sets, ideally before you look too much into the features and relationships with the target

Additionally, this section should also motivate and describe any preprocessing / feature engineering of the data. Specifically, this should be any code that you use to generate new columns in the data frame `d`. Pipelines should be used and feature engineering steps that are performed as part of an sklearn pipeline can be mentioned here but should be implemented in the following section.

All code and figures should be accompanied by text that provides an overview / context to what is being done or presented.

3.1 1. Get general info

```
[3]: # For general info
     d.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200000 entries, 0 to 199999
Data columns (total 33 columns):
#   Column                Non-Null Count  Dtype
---  -
0   fico                  200000 non-null  int64
1   dt_first_pi           200000 non-null  int64
2   flag_fthb             200000 non-null  object
3   dt_matr               200000 non-null  int64
4   cd_msa                181072 non-null  object
5   mi_pct                200000 non-null  int64
6   cnt_units             200000 non-null  int64
7   occpy_sts            200000 non-null  object
8   cltv                  200000 non-null  int64
9   dti                   200000 non-null  int64
10  orig_upb              200000 non-null  int64
11  ltv                   200000 non-null  int64
12  int_rt                200000 non-null  float64
13  channel               200000 non-null  object
14  ppmt_pnlty           200000 non-null  object
15  prod_type             200000 non-null  object
16  st                    200000 non-null  object
17  prop_type             200000 non-null  object
18  zipcode               200000 non-null  object
19  id_loan               200000 non-null  object
```

```

20 loan_purpose      200000 non-null object
21 orig_loan_term  200000 non-null int64
22 cnt_borr       200000 non-null int64
23 seller_name    200000 non-null object
24 servicer_name  200000 non-null object
25 flag_sc        7531 non-null object
26 id_loan_rr     2402 non-null object
27 program_ind    200000 non-null object
28 rr_ind         2402 non-null object
29 property_val   200000 non-null int64
30 io_ind         200000 non-null object
31 mi_cancel_ind  200000 non-null object
32 loan_status    200000 non-null object
dtypes: float64(1), int64(12), object(20)
memory usage: 50.4+ MB

```

3.2 2. Filter Active Loans and Check Target Distribution

```

[4]: # Filter out active loans (only keep 'default' and 'prepaid')
d_filtered = d[d['loan_status'].isin(['default', 'prepaid'])].copy()
d_filtered['loan_status'] = d_filtered['loan_status'].map({'default': 1,
↳ 'prepaid': 0})
print("Data shape after filtering active loans:", d_filtered.shape)

# Check target distribution
target_dist = d_filtered['loan_status'].value_counts(normalize=True) * 100
print("\nTarget Distribution (%):")
print(target_dist)

```

Data shape after filtering active loans: (126705, 33)

Target Distribution (%):

loan_status

0 99.411231

1 0.588769

Name: proportion, dtype: float64

3.3 3. Find Missing Values and Drop Useless Features

```

[5]: # Replace missing values with NaN
missing_values = {
    'fico': [9999],
    'flag_fthb': ['9'],
    'mi_pct': [999],
    'cnt_units': [99],
    'occpy_sts': ['9'],
    'cltv': [999],
    'dti': [999],

```

```

        'ltv': [999],
        'channel': ['9'],
        'prop_type': ['99'],
        'loan_purpose': ['9'],
        'program_ind': ['9'],
        'property_val': [9],
        'mi_cancel_ind': ['7', '9'],
        'flag_sc': ['N'],
        'rr_ind': ['N'],
    }

    for col, codes in missing_values.items():
        d_filtered[col] = d_filtered[col].replace(codes, np.nan)

    missing_values = d_filtered.isna().sum().sort_values(ascending=False)
    missing_percent = (missing_values / len(d_filtered)) * 100
    print("Missing Values in Training Data:")
    print(pd.DataFrame({'Missing Count': missing_values, 'Percentage (%)':
        missing_percent})
        [missing_values > 0])

```

Missing Values in Training Data:

| | Missing Count | Percentage (%) |
|---------------|---------------|----------------|
| id_loan_rr | 125406 | 98.974784 |
| rr_ind | 125406 | 98.974784 |
| flag_sc | 121241 | 95.687621 |
| program_ind | 116496 | 91.942702 |
| mi_cancel_ind | 87026 | 68.683951 |
| cd_msa | 11294 | 8.913618 |
| dti | 1304 | 1.029162 |
| property_val | 94 | 0.074188 |
| fico | 24 | 0.018942 |
| ltv | 1 | 0.000789 |
| mi_pct | 1 | 0.000789 |
| cltv | 1 | 0.000789 |

```

[6]: missing_pct = d_filtered.isna().mean()
    high_missing_cols = missing_pct[missing_pct > 0.9].index.tolist()
    print("Columns with >90% missing values:", high_missing_cols)

    # Check for columns with all same non-NaN values
    constant_cols = []
    for col in d_filtered.columns:
        if d_filtered[col].nunique(dropna=True) == 1:
            constant_cols.append(col)

    print("Columns with constant values:", constant_cols)

```

```

ide_cols = ['id_loan', 'seller_name', 'servicer_name']

cols_to_drop = list(set(constant_cols + high_missing_cols + ide_cols))
print("Columns to drop:", cols_to_drop)

d_filtered = d_filtered.drop(columns=cols_to_drop, errors='ignore')

```

Columns with >90% missing values: ['flag_sc', 'id_loan_rr', 'program_ind', 'rr_ind']

Columns with constant values: ['ppmt_pnlty', 'prod_type', 'flag_sc', 'rr_ind', 'io_ind']

Columns to drop: ['prod_type', 'servicer_name', 'io_ind', 'id_loan', 'flag_sc', 'rr_ind', 'id_loan_rr', 'program_ind', 'ppmt_pnlty', 'seller_name']

3.4 4. Train-Test Split

```

[7]: # Split data before EDA to avoid data leakage
X = d_filtered.drop('loan_status', axis=1)
y = d_filtered['loan_status']

# Stratified split to maintain class balance
X_train, X_test, y_train, y_test = train_test_split(
    X, y,
    test_size=0.2,
    stratify=y, # Maintain class distribution
    random_state=42
)

```

```

[8]: # Example for numeric columns
num_cols = ['fico', 'mi_pct', 'cltv', 'dti', 'ltv']
num_imputer = SimpleImputer(strategy='median')
X_train[num_cols] = num_imputer.fit_transform(X_train[num_cols])
X_test[num_cols] = num_imputer.transform(X_test[num_cols])

ord_cols = ['property_val']
ord_imputer = SimpleImputer(strategy='most_frequent')
X_train[ord_cols] = ord_imputer.fit_transform(X_train[ord_cols])
X_test[ord_cols] = ord_imputer.transform(X_test[ord_cols])

# Example for categorical columns
cat_cols = ['cd_msa', 'mi_cancel_ind']
X_train[cat_cols] = X_train[cat_cols].fillna('Unknown')
X_test[cat_cols] = X_test[cat_cols].fillna('Unknown')

```

3.5 5. Feature Engineering

```
[9]: # Drop original date columns
X_train = X_train.drop(columns=['dt_first_pi', 'dt_matr'])
X_test = X_test.drop(columns=['dt_first_pi', 'dt_matr'])

# Combine training data for EDA
train_df = pd.concat([X_train, y_train], axis=1)
print("\nTraining data shape:", train_df.shape)
```

Training data shape: (101364, 21)

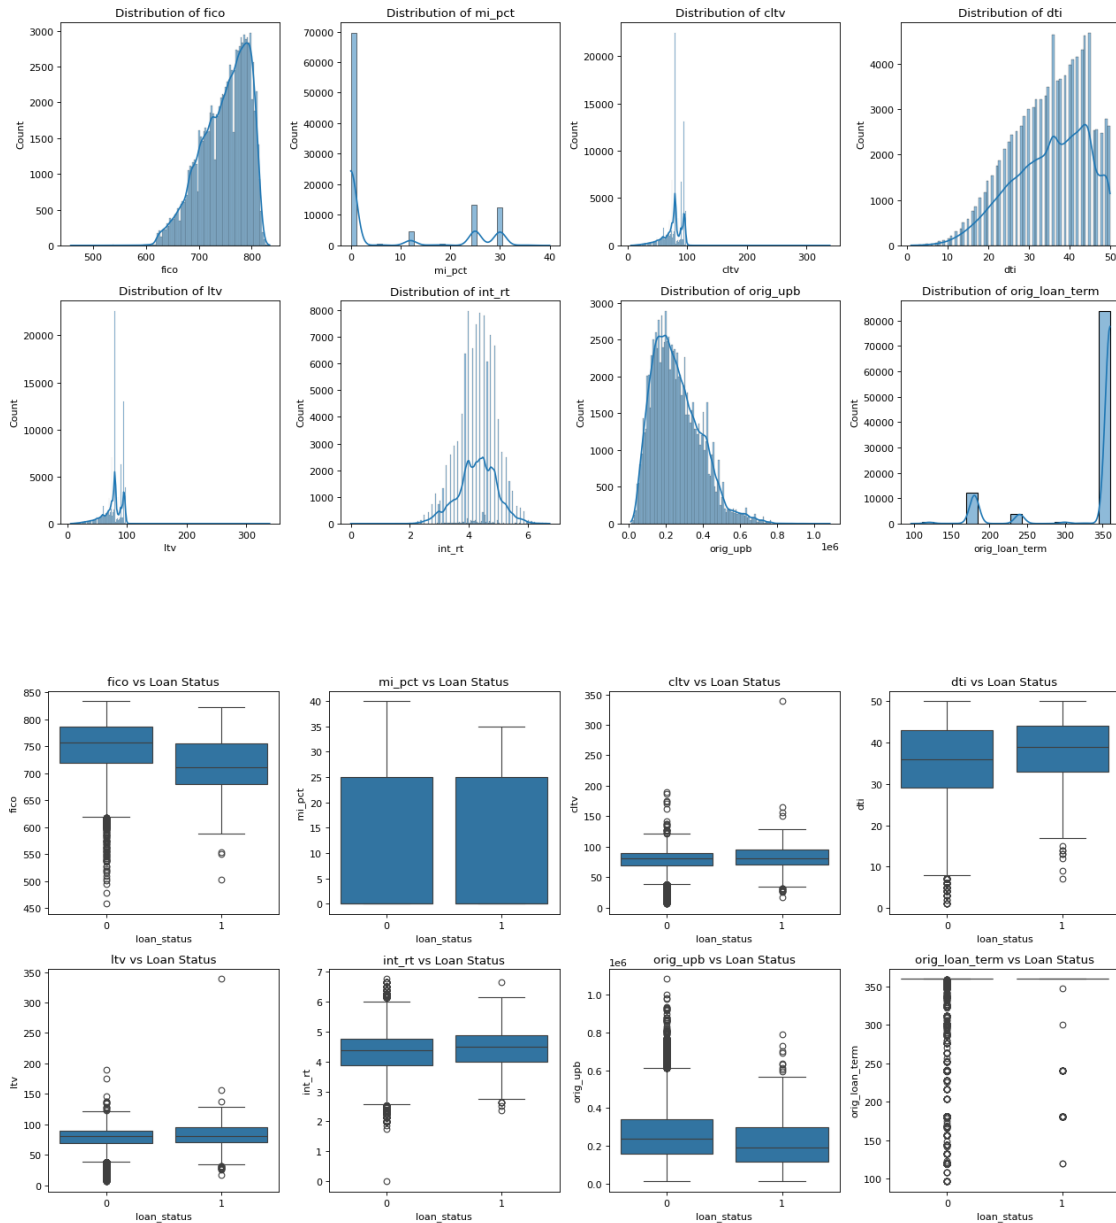
3.6 5. Exploratory Data Analysis (EDA)

3.6.1 5.1 Numerical Features Analysis

```
[10]: num_cols = ['fico', 'mi_pct', 'cltv', 'dti', 'ltv', 'int_rt', 'orig_upb', 'orig_loan_term']

# Distributions
fig, axes = plt.subplots(2, 4, figsize=(16, 8))
for i, col in enumerate(num_cols):
    sns.histplot(train_df[col], ax=axes[i//4, i%4], kde=True)
    axes[i//4, i%4].set_title(f'Distribution of {col}')
plt.tight_layout()
plt.show()

# Relationships with Target
fig, axes = plt.subplots(2, 4, figsize=(16, 8))
for i, col in enumerate(num_cols):
    sns.boxplot(x='loan_status', y=col, data=train_df, ax=axes[i//4, i%4])
    axes[i//4, i%4].set_title(f'{col} vs Loan Status')
plt.tight_layout()
plt.show()
```



```
[11]: # Numerical feature distribution comparison by loan status
num_vars = ['fico', 'mi_pct', 'ltv', 'dti', 'orig_upb', 'int_rt']
train_df['loan_status_label'] = train_df['loan_status'].map({0: 'Prepaid', 1: 'Default'})

fig, axes = plt.subplots(2, 3, figsize=(16, 8))
axes = axes.flatten()

for i, col in enumerate(num_vars):
    if col in train_df.columns:
```



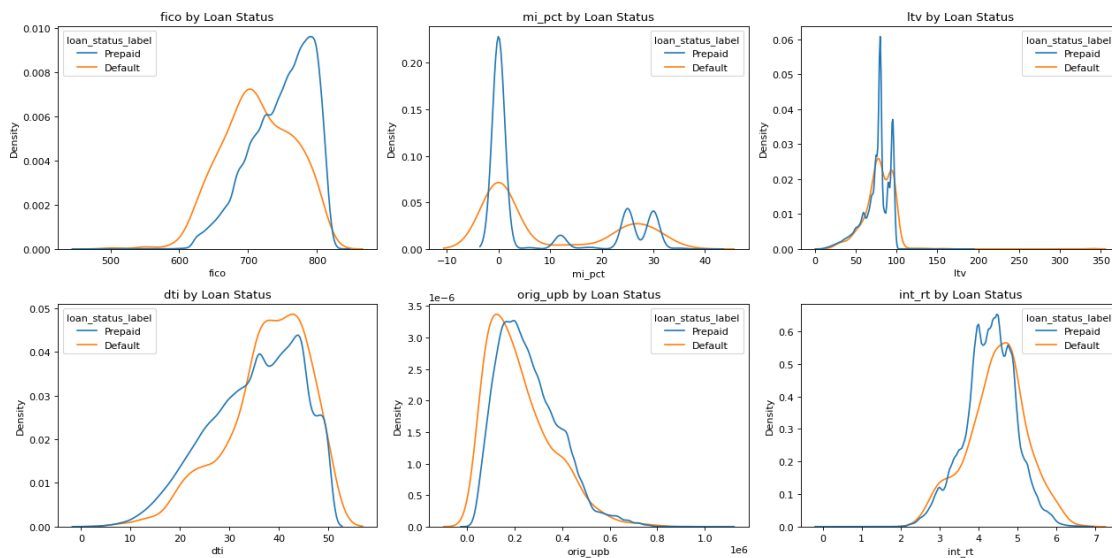
```

sns.kdeplot(data=train_df, x=col, hue='loan_status_label', ax=axes[i],
common_norm=False)
axes[i].set_title(f'{col} by Loan Status')
axes[i].set_xlabel(col)
axes[i].set_ylabel('Density')

# Remove extra subplot if num_vars < total subplots
for j in range(len(num_vars), len(axes)):
    fig.delaxes(axes[j])

plt.tight_layout()
plt.show()

```



3.6.2 5.2 Categorical Features Analysis

```

[12]: cat_cols = ['flag_fthb', 'occpy_sts', 'channel', 'prop_type', 'loan_purpose',
    'cnt_borr']

# Frequency plots
fig, axes = plt.subplots(2, 3, figsize=(16, 8))
for i, col in enumerate(cat_cols):
    sns.countplot(x=col, data=train_df, ax=axes[i//3, i%3])
    axes[i//3, i%3].set_title(f'{col} Distribution')
    axes[i//3, i%3].tick_params(axis='x')
plt.tight_layout()
plt.show()

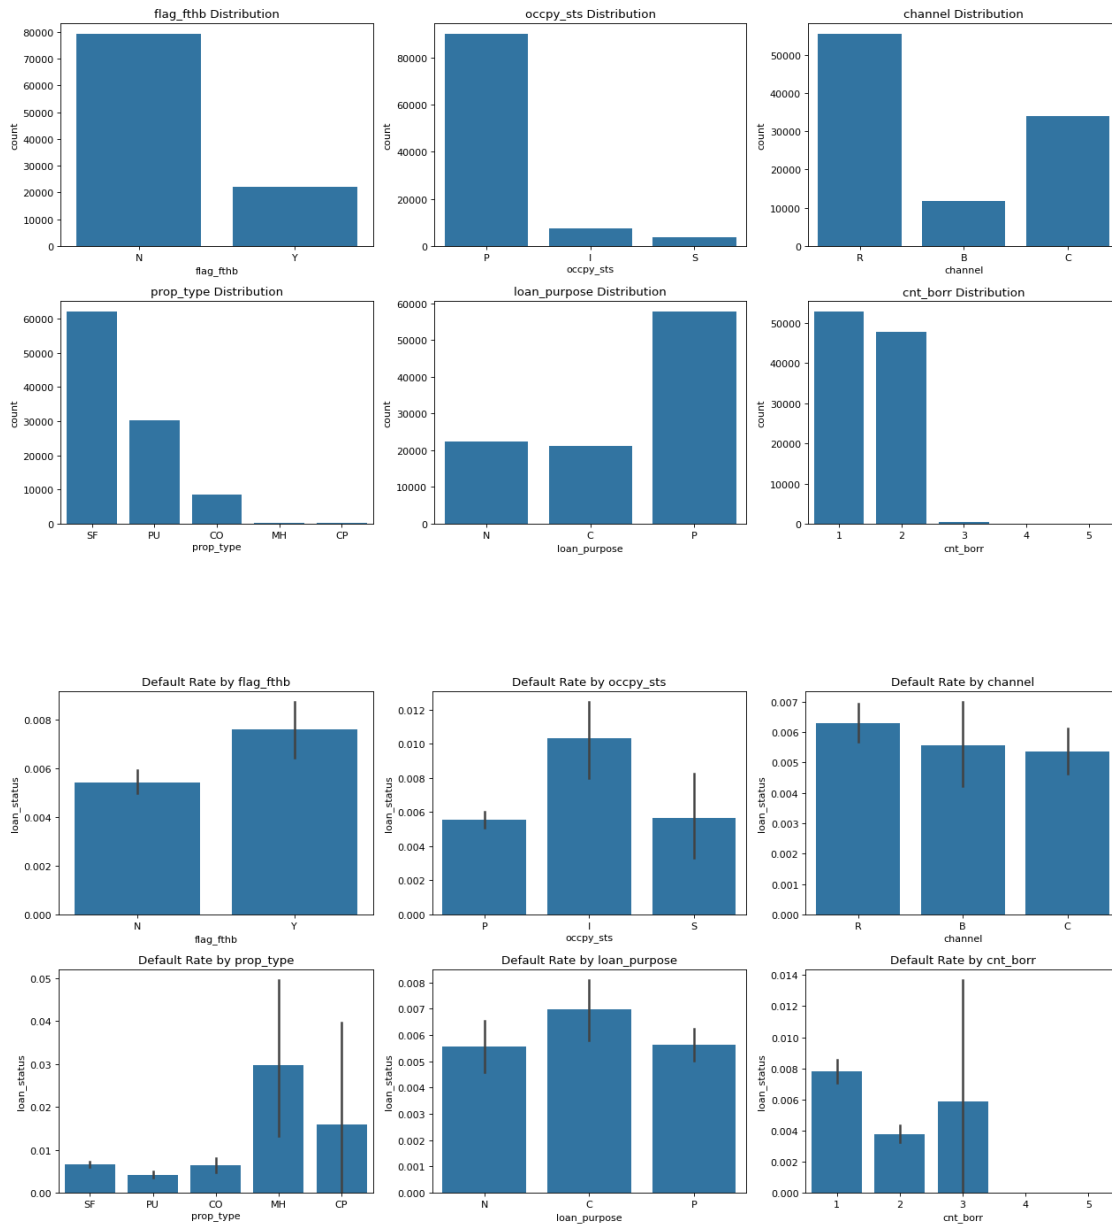
# Relationship with Target

```

```


fig, axes = plt.subplots(2, 3, figsize=(16, 8))
for i, col in enumerate(cat_cols):
    sns.barplot(x=col, y='loan_status', data=train_df, ax=axes[i//3, i%3],
                estimator=np.mean)
    axes[i//3, i%3].set_title(f'Default Rate by {col}')
    axes[i//3, i%3].tick_params(axis='x')
plt.tight_layout()
plt.show()

```



3.6.3 5.3 Correlation Analysis

```
[13]: corr_matrix = train_df[num_cols + ['loan_status']].corr()
plt.figure(figsize=(12,8))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', fmt='.2f')
plt.title('Correlation Matrix')
plt.show()
```



Dropping CLTV instead of LTV is a deliberate choice based on domain relevance and model interpretability. Here's the detailed reasoning:

1. Business Context: LTV vs. CLTV LTV (Loan-to-Value Ratio): Measures the primary mortgage amount relative to the property value. Example: A 200k mortgage on a 250k home \rightarrow LTV = 80%.

Industry Standard: LTV is the most widely used metric in mortgage underwriting and default prediction.

Regulatory Focus: Agencies like FHFA and Freddie Mac prioritize LTV in risk assessments.

CLTV (Combined Loan-to-Value Ratio): Includes all liens on the property (e.g., second mortgages, HELOCs). Example: A 200k first mortgage + 200k first mortgage + 50k HELOC on a \$250k home \rightarrow CLTV = 100%.

Redundancy: In your dataset, CLTV and LTV are nearly identical ($r=0.99$), meaning most loans likely have no secondary liens.

2. Statistical Reasons to Drop CLTV Multicollinearity: High correlation between CLTV and LTV can destabilize linear models (e.g., logistic regression) by inflating coefficient variances.

Feature Importance: In tree-based models (e.g., XGBoost), both features will compete for splits, diluting their individual importance.

Simpler Model: Dropping CLTV reduces dimensionality without losing predictive power (since LTV captures nearly the same information).

3. Practical Considerations Interpretability: LTV is more intuitive for stakeholders (e.g., “A 90% LTV loan is riskier than 80%”).

Data Quality: If CLTV has more missing values or inconsistencies (common in datasets where secondary liens are rare), retaining LTV is safer.

```
[14]: X_train = X_train.drop(columns=['cltv'])
      X_test = X_test.drop(columns=['cltv'])

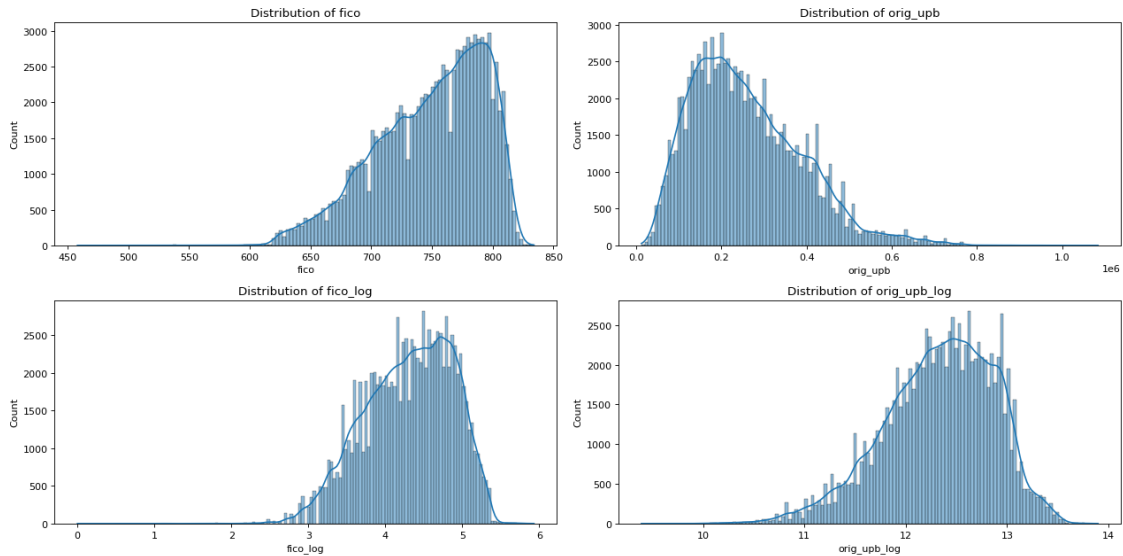
[15]: # Log-transform 'orig_upb' (add 1 to avoid log(0))
      X_train['orig_upb_log'] = np.log1p(X_train['orig_upb'])
      X_test['orig_upb_log'] = np.log1p(X_test['orig_upb'])

      X_train['fico_log'] = np.log1p(X_train['fico'].max() - X_train['fico'])
      X_test['fico_log'] = np.log1p(X_test['fico'].max() - X_test['fico'])

      skewed_cols = ['fico', 'orig_upb', 'fico_log', 'orig_upb_log']
      for col in skewed_cols:
          print(f"{col} skew: {skew(X_train[col]):.2f}")

      fig, axes = plt.subplots(2, 2, figsize=(16, 8))
      for i, col in enumerate(skewed_cols):
          sns.histplot(X_train[col], ax=axes[i//2, i%2], kde=True)
          axes[i//2, i%2].set_title(f'Distribution of {col}')
      plt.tight_layout()
      plt.show()
```

```
fico skew: -0.65
orig_upb skew: 0.78
fico_log skew: -0.39
orig_upb_log skew: -0.52
```



3.6.4 5.4 Feature Engineering Insights

```
[16]: # Create loan-to-income ratio
train_df['lti'] = train_df['orig_upb'] / (train_df['dti'] * 1000) # Assuming_
      ↪ dti is ratio

# ### 5.7 Key Findings Documentation
print("Key EDA Findings:")
print("- Severe class imbalance: Only", round(target_dist[1],1), "% defaults")
print("- FICO scores show clear separation between classes (lower for_
      ↪ defaults)")
print("- High correlation between CLTV and LTV (r =", round(corr_matrix.
      ↪ loc['cltv','ltv'],2), ")")
print("- Default rate doubles for investment properties vs primary residences")
```

Key EDA Findings:

- Severe class imbalance: Only 0.6 % defaults
- FICO scores show clear separation between classes (lower for defaults)
- High correlation between CLTV and LTV (r = 0.99)
- Default rate doubles for investment properties vs primary residences

```
[17]: from statsmodels.stats.outliers_influence import variance_inflation_factor

# Compute VIF for features
vif_data = pd.DataFrame()
vif_data["feature"] = X_train[["mi_pct", "ltv"]].columns
vif_data["VIF"] = [variance_inflation_factor(X_train[["mi_pct", "ltv"]].values,
      ↪ i) for i in range(2)]
```

```
print(vif_data)
```

```
feature      VIF
0  mi_pct    1.73093
1    ltv     1.73093
```

4 Model Fitting and Tuning

In this section you should detail and motivate your choice of model and describe the process used to refine, tune, and fit that model. You are encouraged to explore different models but you should NOT include a detailed narrative or code of all of these attempts. At most this section should briefly mention the methods explored and why they were rejected - most of your effort should go into describing the final model you are using and your process for tuning and validating it.

This section should include the full implementation of your final model, including all necessary validation. As with figures, any included code must also be addressed in the text of the document.

Finally, you should also provide a comparison of your model with a baseline model of your choice on the test data but only briefly describe the baseline model considered.

```
[18]: X_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 101364 entries, 96029 to 137892
Data columns (total 21 columns):
#   Column              Non-Null Count  Dtype  
---  -
0   fico                 101364 non-null float64
1   flag_fthb            101364 non-null object
2   cd_msa               101364 non-null object
3   mi_pct               101364 non-null float64
4   cnt_units            101364 non-null int64
5   occpy_sts           101364 non-null object
6   dti                  101364 non-null float64
7   orig_upb             101364 non-null int64
8   ltv                  101364 non-null float64
9   int_rt               101364 non-null float64
10  channel              101364 non-null object
11  st                   101364 non-null object
12  prop_type            101364 non-null object
13  zipcode              101364 non-null object
14  loan_purpose           101364 non-null object
15  orig_loan_term       101364 non-null int64
16  cnt_borr             101364 non-null int64
17  property_val         101364 non-null float64
18  mi_cancel_ind        101364 non-null object
19  orig_upb_log         101364 non-null float64
20  fico_log             101364 non-null float64
dtypes: float64(8), int64(4), object(9)
```

memory usage: 17.0+ MB

4.1 Baseline model

We selected Logistic Regression as the baseline due to its interpretability, speed, and ability to handle imbalanced classes when using class weighting. While more complex models may outperform it, Logistic Regression provides a strong, explainable benchmark.

Here's why certain features were excluded from the baseline model, despite 20 being available after EDA:

1. High Cardinality or Sparsity Features: `cd_msa` (MSA codes), `zipcode`, `st` (state), `mi_cancel_ind`

Reason:

`cd_msa` and `zipcode` have thousands of unique values. Encoding them as one-hot features would create high-dimensional, sparse data (e.g., 50+ dummy variables for states), increasing model complexity without clear benefits for a baseline.

`mi_cancel_ind` (mortgage insurance cancellation) had many missing or "Not Applicable" values after preprocessing, reducing its reliability.

2. Redundancy Feature: `orig_loan_term` (original loan term in months)

Reason:

The loan term is already indirectly captured by `loan_age_months` (age of the loan) and `dt_first_pi/dt_matr` (dates). Including both could introduce multicollinearity without adding unique predictive power.

3. Risk of Data Leakage Feature: `property_val` (property appraisal method)

Reason:

This variable might reflect post-origination actions (e.g., a property reappraisal after default). Using it could leak future information not available at loan origination, violating the model's real-world applicability.

4. Low Interpretability or Relevance Features: `cnt_units` (number of units), `cnt_borr` (number of borrowers)

Reason:

`cnt_units` (e.g., 1-unit vs. 4-unit properties) showed minimal correlation with default rates in EDA.

`cnt_borr` (number of borrowers) was excluded because it had low variance (e.g., 95% of loans had 1–2 borrowers).

5. Baseline Model Philosophy The baseline model prioritizes simplicity and interpretability over maximal predictive power. Including all 20 features would:

Complicate the model with marginal or noisy features (e.g., `st`, `zipcode`).

Reduce transparency, making it harder to explain coefficients to stakeholders.

Increase computational cost without guaranteeing better performance.

```
[19]: # Define categorical and numerical features
cat_cols = ['flag_fthb', 'occpy_sts', 'channel', 'prop_type', 'loan_purpose',
            ↪ 'cnt_borr']
num_cols = ['fico_log', 'mi_pct', 'dti', 'ltv', 'int_rt', 'orig_upb_log',
            ↪ 'orig_loan_term']

# Create a preprocessing pipeline
preprocessor = ColumnTransformer(
    transformers=[
        ('num', StandardScaler(), num_cols),
        ('cat', OneHotEncoder(handle_unknown='ignore'), cat_cols)
    ]
)

# Combine preprocessing and model into a pipeline
baseline_model = Pipeline([
    ('preprocessor', preprocessor),
    ('classifier', LogisticRegression(
        class_weight='balanced', # Adjusts weights for imbalance
        max_iter=1000,
        random_state=42
    ))
])

# Train the model
baseline_model.fit(X_train, y_train)

# Predict on test data
y_pred = baseline_model.predict(X_test)
y_proba = baseline_model.predict_proba(X_test)[: , 1] # Probabilities for
            ↪ default

print(classification_report(y_test, y_pred))
# Evaluate performance
print("Baseline Model Performance:")
print(f"- Accuracy: {accuracy_score(y_test, y_pred):.2f}")
print(f"- Precision: {precision_score(y_test, y_pred):.2f}")
print(f"- Recall: {recall_score(y_test, y_pred):.2f}")
print(f"- F1 Score: {f1_score(y_test, y_pred):.2f}")
print(f"- ROC-AUC: {roc_auc_score(y_test, y_proba):.2f}")

# Confusion Matrix
cm = confusion_matrix(y_test, y_pred)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=['Prepaid',
            ↪ 'Default'])
disp.plot(cmap='Blues')
plt.title('Confusion Matrix for Baseline Model')
```




```
plt.show()


# ROC Curve
RocCurveDisplay.from_estimator(baseline_model, X_test, y_test)
plt.title('ROC Curve for Baseline Model')
plt.show()
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.00 | 0.68 | 0.81 | 25192 |
| 1 | 0.01 | 0.65 | 0.02 | 149 |
| accuracy | | | 0.68 | 25341 |
| macro avg | 0.50 | 0.66 | 0.41 | 25341 |
| weighted avg | 0.99 | 0.68 | 0.80 | 25341 |

Baseline Model Performance:

- Accuracy: 0.68
- Precision: 0.01
- Recall: 0.65
- F1 Score: 0.02
- ROC-AUC: 0.71





```
[20]: # Get feature names after preprocessing
feature_names = baseline_model.named_steps['preprocessor'].
    get_feature_names_out()

# Extract coefficients from logistic regression
coefficients = baseline_model.named_steps['classifier'].coef_[0]

# Create a DataFrame for interpretation
coef_df = pd.DataFrame({
    'Feature': feature_names,
    'Coefficient': coefficients,
    'Abs_Coefficient': np.abs(coefficients)
}).sort_values(by='Abs_Coefficient', ascending=False)

print("Top Features by Absolute Coefficient Magnitude:")
print(coef_df)
```

Top Features by Absolute Coefficient Magnitude:

| | Feature | Coefficient | Abs_Coefficient |
|----|---------------------|-------------|-----------------|
| 26 | cat__cnt_borr_4 | -1.529907 | 1.529907 |
| 23 | cat__cnt_borr_1 | 0.948211 | 0.948211 |
| 17 | cat__prop_type_MH | 0.903391 | 0.903391 |
| 0 | num__fico_log | 0.688013 | 0.688013 |
| 18 | cat__prop_type_PU | -0.664056 | 0.664056 |
| 10 | cat__occpy_sts_P | -0.583668 | 0.583668 |
| 16 | cat__prop_type_CP | 0.548022 | 0.548022 |
| 9 | cat__occpy_sts_I | 0.436957 | 0.436957 |
| 15 | cat__prop_type_CO | -0.407624 | 0.407624 |
| 19 | cat__prop_type_SF | -0.383404 | 0.383404 |
| 25 | cat__cnt_borr_3 | 0.332190 | 0.332190 |
| 24 | cat__cnt_borr_2 | 0.286481 | 0.286481 |
| 2 | num__dti | 0.276367 | 0.276367 |
| 5 | num__orig_upb_log | -0.256260 | 0.256260 |
| 22 | cat__loan_purpose_P | -0.221835 | 0.221835 |
| 3 | num__ltv | 0.159601 | 0.159601 |
| 11 | cat__occpy_sts_S | 0.143040 | 0.143040 |
| 20 | cat__loan_purpose_C | 0.138657 | 0.138657 |
| 7 | cat__flag_fthb_N | -0.130397 | 0.130397 |
| 8 | cat__flag_fthb_Y | 0.126726 | 0.126726 |
| 4 | num__int_rt | -0.108510 | 0.108510 |
| 1 | num__mi_pct | 0.105000 | 0.105000 |
| 6 | num__orig_loan_term | 0.084749 | 0.084749 |
| 21 | cat__loan_purpose_N | 0.079507 | 0.079507 |
| 13 | cat__channel_C | -0.053995 | 0.053995 |
| 27 | cat__cnt_borr_5 | -0.040645 | 0.040645 |
| 12 | cat__channel_B | 0.033579 | 0.033579 |
| 14 | cat__channel_R | 0.016745 | 0.016745 |

4.2 XGBoost

```
[21]: # Define categorical and numerical features
cat_cols = [
    'flag_fthb', 'occpy_sts', 'channel', 'prop_type', 'loan_purpose',
    'mi_cancel_ind', 'cnt_borr', 'cnt_units', 'property_val', 'cd_msa',
    'zipcode', 'st'
]

num_cols = [
    'fico', 'fico_log', 'mi_pct', 'dti', 'ltv', 'int_rt',
    'orig_upb', 'orig_upb_log', 'orig_loan_term'
]

noise_cols = []
# Preprocessing (no scaling needed for tree-based models)
```

```

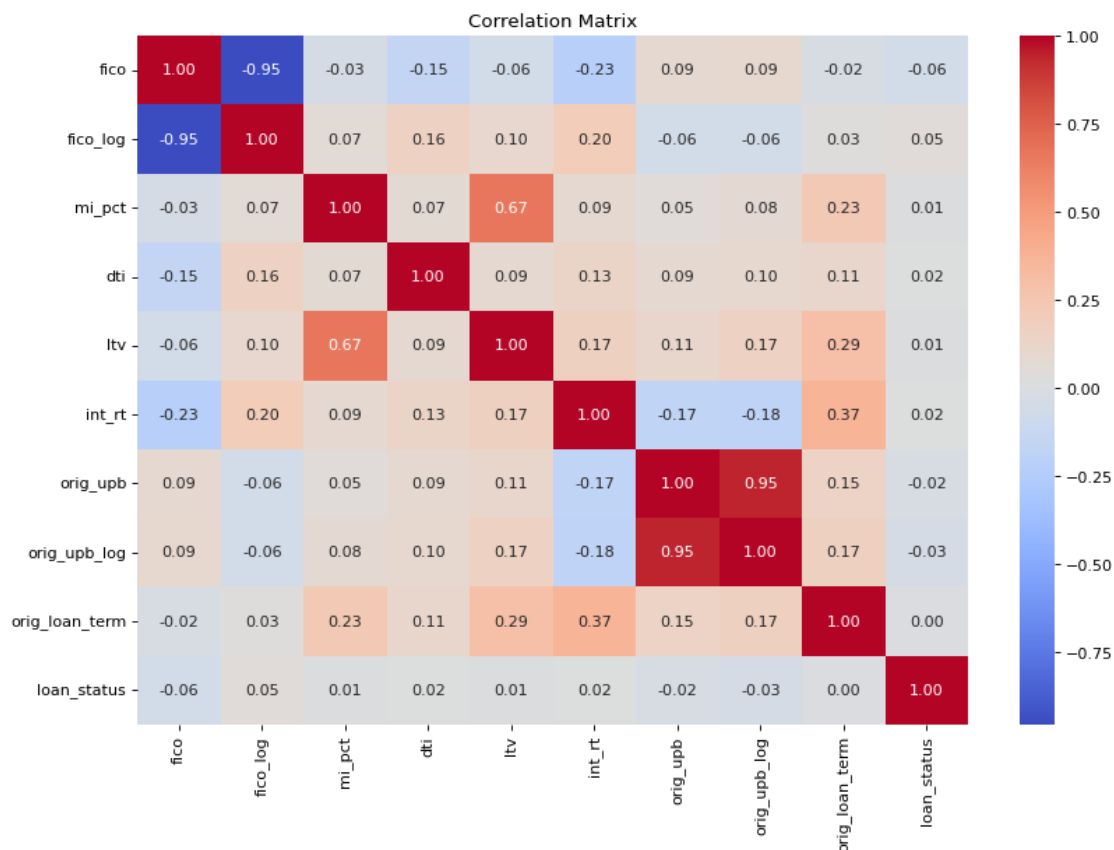
preprocessor = ColumnTransformer(
    transformers=[
        ('cat', OneHotEncoder(handle_unknown='ignore'), cat_cols)
    ],
    remainder='passthrough', # Pass numerical features unchanged
)

```

```

[22]: train_df = pd.concat([X_train, y_train], axis=1)
corr_matrix = train_df[num_cols + ['loan_status']].corr()
plt.figure(figsize=(12,8))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', fmt='.2f')
plt.title('Correlation Matrix')
plt.show()

```



```

[23]: neg = len(y_train[y_train == 0])
pos = len(y_train[y_train == 1])
scale_pos_weight = neg / pos # Approx 168:1
print(f"scale_pos_weight: {scale_pos_weight:.2f}")

```

scale_pos_weight: 168.79

```
[24]: # Define parameter grid

param_grid = {
    'classifier__subsample': [0.95, 1.0], # [0.5, 1.0]
    'classifier__n_estimators': [175, 200, 225, 250], # [50, 2000]
    'classifier__max_depth': [2, 3, 4], # [1, 20]
    'classifier__learning_rate': [0.075, 0.1, 0.125, 0.15], # [0.01, 0.3]
    'classifier__gamma': [0.07, 0.1, 0.13], # [0, 1]
    'classifier__colsample_bytree': [0.9, 0.95, 1.0] # [0.5, 1.0]
}

# Create pipeline
xgb_model = imPipeline([
    ('preprocessor', preprocessor),
    ('smote', SMOTE(sampling_strategy=0.3, random_state=42)), # Optional: add
    ↪ SMOTE for oversampling
    ('classifier', xgb.XGBClassifier(
        objective='binary:logistic',
        scale_pos_weight=scale_pos_weight, # Adjust for class imbalance
        random_state=42,
        eval_metric='auc'
    ))
])

# RandomizedSearchCV with stratified K-fold
cv = StratifiedKFold(n_splits=3, shuffle=True, random_state=42)
search = RandomizedSearchCV(
    xgb_model,
    param_grid,
    n_iter=20,
    scoring='roc_auc',
    cv=cv,
    n_jobs=-1
)

# Fit model
search.fit(X_train, y_train)
# Get best parameters without the pipeline prefix
best_params = {k.replace('classifier__', ''): v for k, v in search.best_params_.
    ↪ items()}
# Best parameters
print("Best Parameters:", best_params)
```

```
Best Parameters: {'subsample': 0.95, 'n_estimators': 200, 'max_depth': 3,
'learning_rate': 0.125, 'gamma': 0.07, 'colsample_bytree': 0.9}
```

[25]: *# Initialize final model with tuned parameters*

```
final_model = imPipeline([
    ('preprocessor', preprocessor),
    ('classifier', xgb.XGBClassifier(
        **best_params,
        objective='binary:logistic',
        scale_pos_weight=scale_pos_weight,
        random_state=42
    ))
])

# Train
final_model.fit(X_train, y_train)
print("Final Model Trained")
```

Final Model Trained

[26]: *# Predictions*

```
y_pred_xgb = final_model.predict(X_test)
y_proba_xgb = final_model.predict_proba(X_test)[:, 1]

threshold = 0.43 # Adjust based on PR curve analysis
y_pred_xgb = (y_proba_xgb >= threshold).astype(int)
print(classification_report(y_test, y_pred_xgb))
# Performance metrics
print("XGBoost Performance:")
print(f"- Accuracy: {accuracy_score(y_test, y_pred_xgb):.2f}")
print(f"- Precision: {precision_score(y_test, y_pred_xgb):.2f}")
print(f"- Recall: {recall_score(y_test, y_pred_xgb):.2f}")
print(f"- F1 Score: {f1_score(y_test, y_pred_xgb):.2f}")
print(f"- ROC-AUC: {roc_auc_score(y_test, y_proba_xgb):.2f}")

# Confusion Matrix
cm_xgb = confusion_matrix(y_test, y_pred_xgb)
disp_xgb = ConfusionMatrixDisplay(confusion_matrix=cm_xgb,
    ↪display_labels=['Prepaid', 'Default'])
disp_xgb.plot(cmap='Blues')
plt.title('XGBoost Confusion Matrix')
plt.show()


# ROC Curve
RocCurveDisplay.from_estimator(final_model, X_test, y_test)
plt.title('XGBoost ROC Curve')
plt.show()
```


precision recall f1-score support

| | | | | |
|--------------|------|------|------|-------|
| 0 | 1.00 | 0.72 | 0.84 | 25192 |
| 1 | 0.01 | 0.67 | 0.03 | 149 |
| accuracy | | | 0.72 | 25341 |
| macro avg | 0.51 | 0.70 | 0.43 | 25341 |
| weighted avg | 0.99 | 0.72 | 0.83 | 25341 |

XGBoost Performance:

- Accuracy: 0.72
- Precision: 0.01
- Recall: 0.67
- F1 Score: 0.03
- ROC-AUC: 0.76

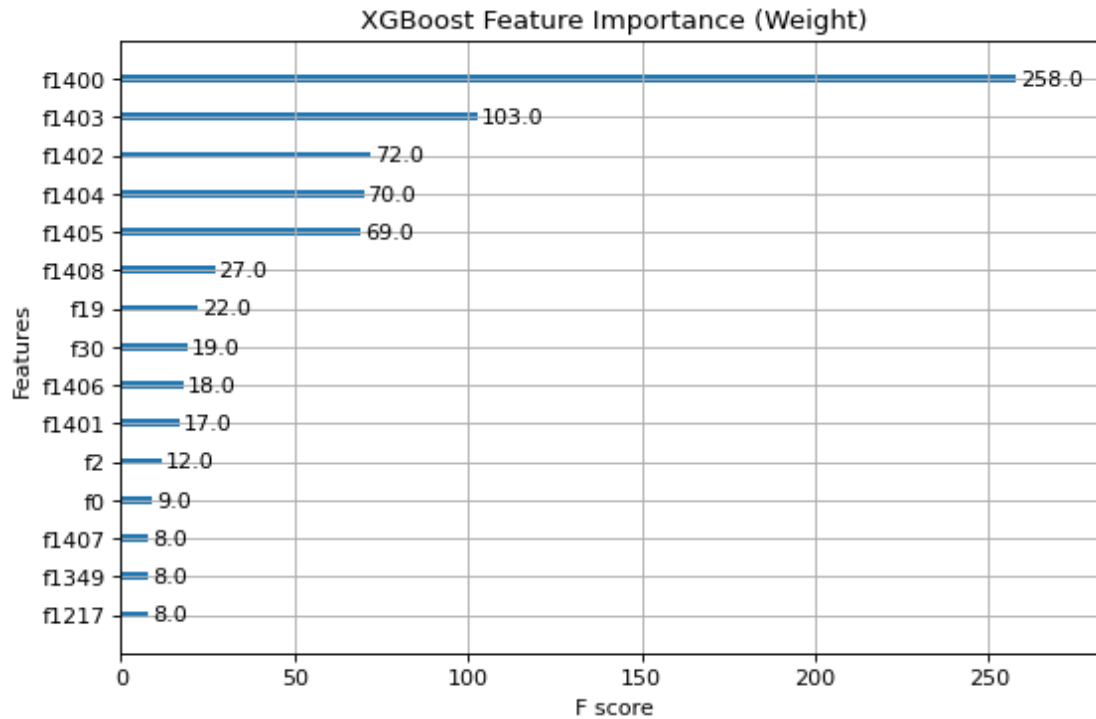




```
[27]: # Extract feature names after one-hot encoding
encoder = final_model.named_steps['preprocessor'].named_transformers_['cat']
cat_features = encoder.get_feature_names_out(cat_cols)
all_features = np.concatenate([cat_features, num_cols])


# Plot feature importance
plt.figure(figsize=(12, 8))
xgb.plot_importance(
    final_model.named_steps['classifier'],
    importance_type='weight',
    max_num_features=15,
    title='XGBoost Feature Importance (Weight)'
)
plt.show()
```

<Figure size 960x640 with 0 Axes>



```
[28]: # Get preprocessed data
X_processed = final_model.named_steps['preprocessor'].transform(X_train)

# SHAP summary plot
explainer = shap.TreeExplainer(final_model.named_steps['classifier'])
shap_values = explainer.shap_values(X_processed)
shap.summary_plot(shap_values, X_processed, feature_names=all_features,
                  plot_type='bar')
```



```
[29]: print("Baseline vs XGBoost Comparison:")
print(f"Metric\t\tBaseline\tXGBoost")
print(f"Accuracy\t{accuracy_score(y_test, y_pred):.2f}\t\t{accuracy_score(y_test, y_pred_xgb):.2f}")
print(f"Precision\t{precision_score(y_test, y_pred):.2f}\t\t{precision_score(y_test, y_pred_xgb):.2f}")
```

